

MySQL Concrete Architecture (Storage Management)

November 1, 2017

Team

Ante Pimentel

Dusan Birtasevic

- Francis Okoyo
- Hashim Al-Helli

Richard Van

Introduction

- This presentation illustrates the divergences between the conceptual and concrete architecture of a system.
- Our topic of interest is the storage management component of the MySQL Server.
- The learning outcome of this presentation is a better understanding of why certain differences exist between architectures.

Layout

- Subsystems & Interactions

Derivation Process

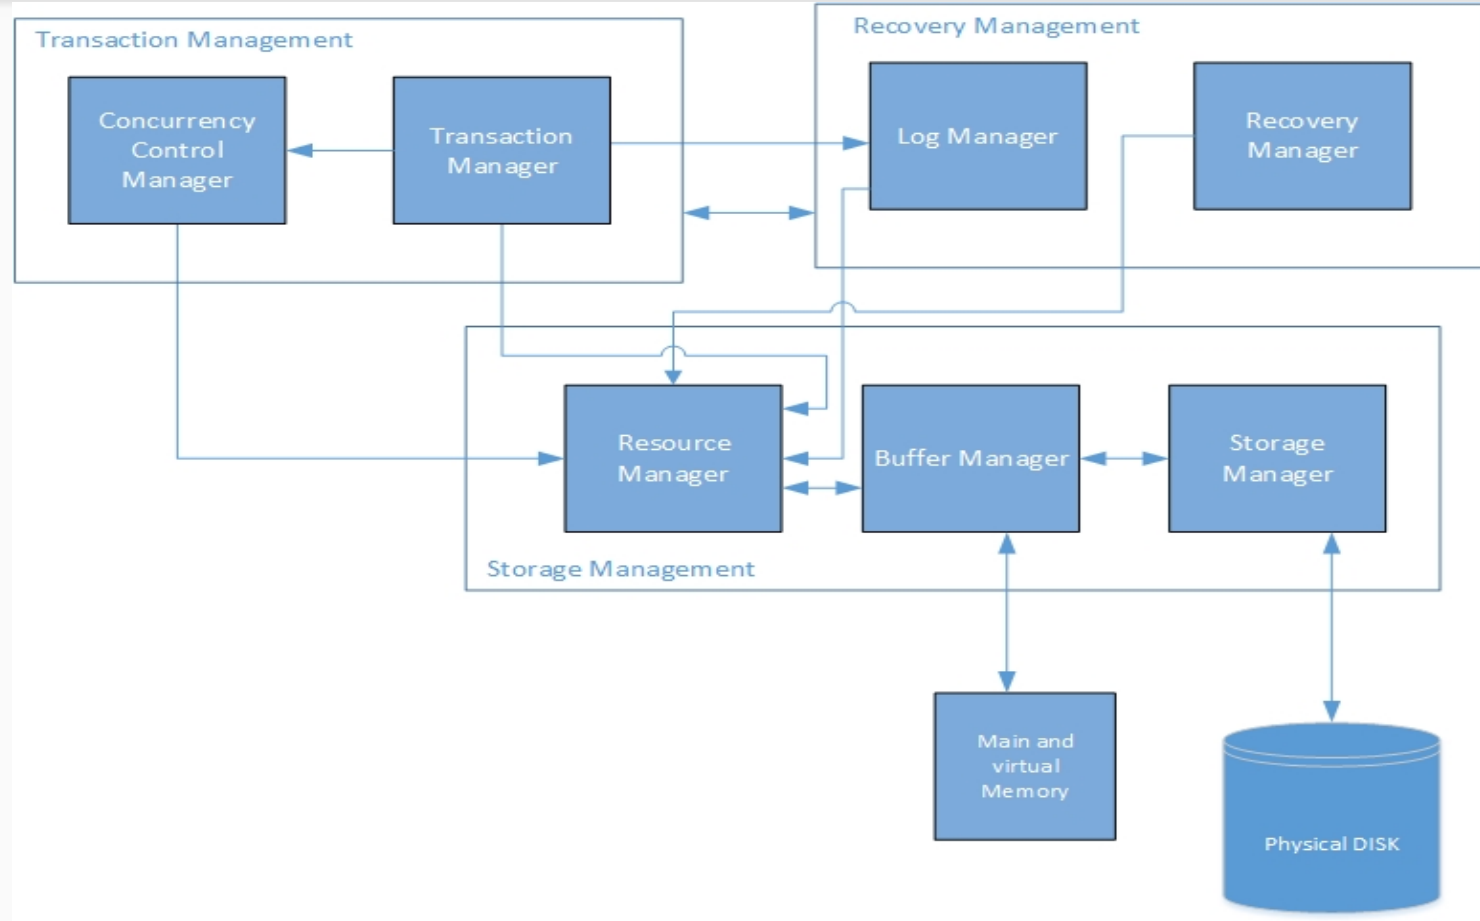
- Architecture Styles & Design Patterns

Storage Management

- Architecture Differences
- LsEdit Screenshots
- Use Case

Lessons Learned

Subsystems and Interactions



Derivation Process

- Visualized interactions using Isedit.
- Simplified the view of Isedit by removing unneeded file dependences.
- Analyzed the code structure and file names and tried to map file names to components.
- If filename wasn't clear we went inside the code and analyzed the comments.
- When we found the right files, we contained and visualized them.

Some Important files in MySQL Server Folder

client - command line client utility code.

sql - main folder of MySQL it contains the parser and optimizer, execution engine and handler to choose a storage engine and many other things.

Sql-common - files that used by client and server.

Storage - The storage engines like **InnoDB** and **MyISAM** implemented here.

Vio - virtual input output contains implementation of low level network I/O code.

Architecture Styles and Design Patterns

- Layered architecture (top-level and Storage management internals)
- Pipe-and-filter architecture (query processing)
- Client-server architecture (application layer)
- Facade design pattern. The file `mysqld.cc` act as an interface between the client and the other components.

Storage Management - Storage Manager

Storage Manager

- Responsible for providing fast read/write access to tablespaces and logs of the database.
- `fil0fil.cc` located in the `fil` directory contains the implementation of the low-level file system. Contains the methods and structures needed to manipulate them.
- `fsp0fsp.cc` located in the `fsp` directory takes care of the file space management, which handles the allocation of pages and keeps track of which files are used, open, and closed.

Storage Management - Buffer Manager

Buffer Manager

- The Buffer Manager is responsible for efficiently storing the data in memory for manipulation
- it accepts requests from the Resource Manager and decides how much memory to allocate.
- The buffer manager is located in the buf directory.
- buf0buf.cc file is the buffer pool, buf0flu.cc flush the buffer pool and buf0lru.cc contains the replacement algorithm that decides which blocks should be shifted back to disk.
- buf0rea.cc calls the storage manager to initiate a file read.

Storage Management - Resource Manager

Resource Manager

- The responsibility of the Resource Manager is to accept requests from the higher level layers and translate them into an appropriate format that can be understood by the Buffer Manager.
- The resource manager is found inside the dict directory.
- It can be seen from the diagram that there is a direct interaction between storage manager and resource manager where the resource manager skips the buffer and accesses the storage manager directly.

Architecture Differences: Sticky Notes

Which: dict0dict.cc (ResourceManager) -> fil0fil.cc (StorageManager)
fil_make_filepath

Who: Kevin Lewis

When: Jan 17, 2014

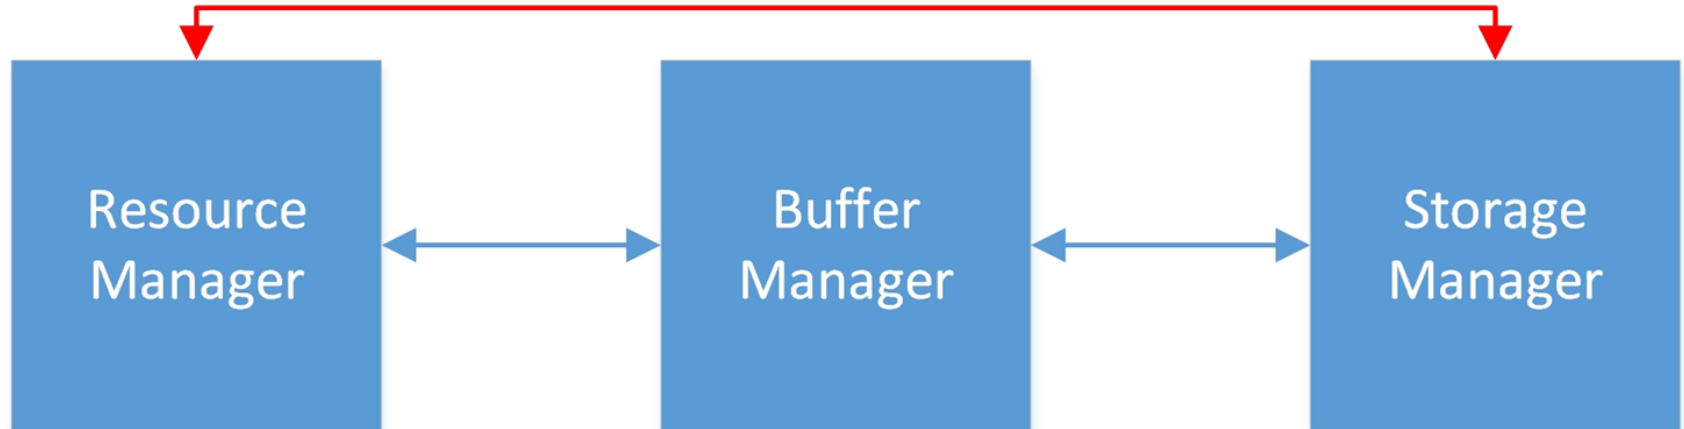
Why: Refactoring. Consolidating several functions that make a filename.

Legend

→ Convergence

→ Divergence

Storage Management



3750	3751		
3751	-		filepath = os_file_make_remote_pathname(
3752	-		table->data_dir_path, table->name, "ibd");
3752	+		filepath = fil_make_filepath(
3753	+		table->data_dir_path,
3754	+		table->name, IBD, true);
3753	3755		} else if (table->dir_path_of_temp_table) {
3754	-		filepath = fil_make_ibd_name(
3755	-		table->dir_path_of_temp_table, true);
3756	+		filepath = fil_make_filepath(
3757	+		table->dir_path_of_temp_table,
3758	+		NULL, IBD, false);
3756	3759		} else {
3757	-		filepath = fil_make_ibd_name(tablename, false);
3760	+		filepath = fil_make_filepath(
3761	+		NULL, tablename, IBD, false);
3759	3761		,

Architecture Differences: Sticky Notes

Which: fil0fil.cc (StorageManager) -> dict0mem.cc (ResourceManager)
dict_mem_create_temporary_tablename

Who: Jimmy Yang

When: Sept 5, 2012

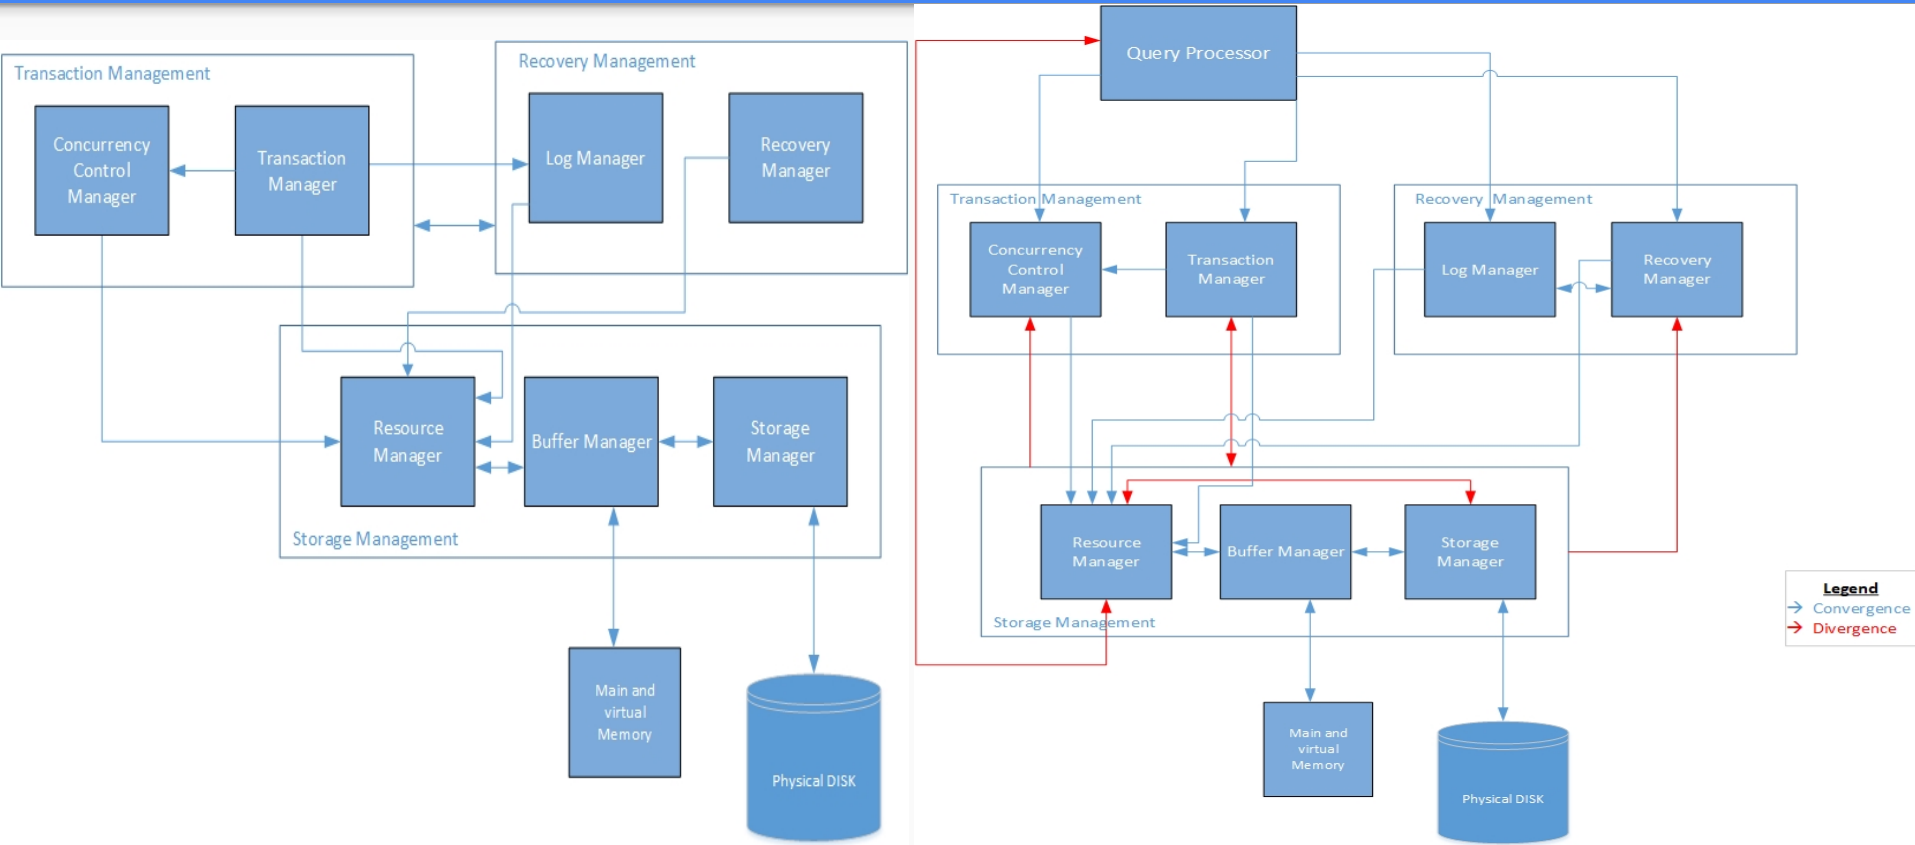
Why: Bug fixing

```


604 +dict_mem_create_temporary_tablename(
605 +/*=====*/
606 +    mem_heap_t*    heap,    /*!< in: memory heap */
607 +    const char*    dbtab,   /*!< in: database/table name */
608 +    table_id_t     id)      /*!< in: InnoDB table id */
609 +{
610 +    const char*     dbend    = strchr(dbtab, '/');
611 +    ut_ad(dbend);
612 +    size_t          dblen    = dbend - dbtab + 1;
613 +    size_t          size     = tmp_file_prefix_length + 4 + 9 + 9 + dblen;
614 +
615 +    char*   name = static_cast<char*>(mem_heap_alloc(heap, size));
616 +    memcpy(name, dbtab, dblen);
617 +    ut_snprintf(name + dblen, size - dblen,
618 +               tmp_file_prefix "-ib" UINT64PF, id);
619 +    return(name);
620 +}

```


Side by side comparison: Conceptual and Concrete architectures




LsEdit Screenshots





mysql-server-mysql-8.0.2/mysql-test/lib/My/SafeProcess .
mysql-server-mysql-8.0.2/mysql-test/lib/My/SafeProcess/safe_process.cc


mysql-server-mysql-8.0.2/scripts .
mysql-server-mysql-8.0.2/scripts/comp_sql.c

mysql-server-mysql-8.0.2/sql

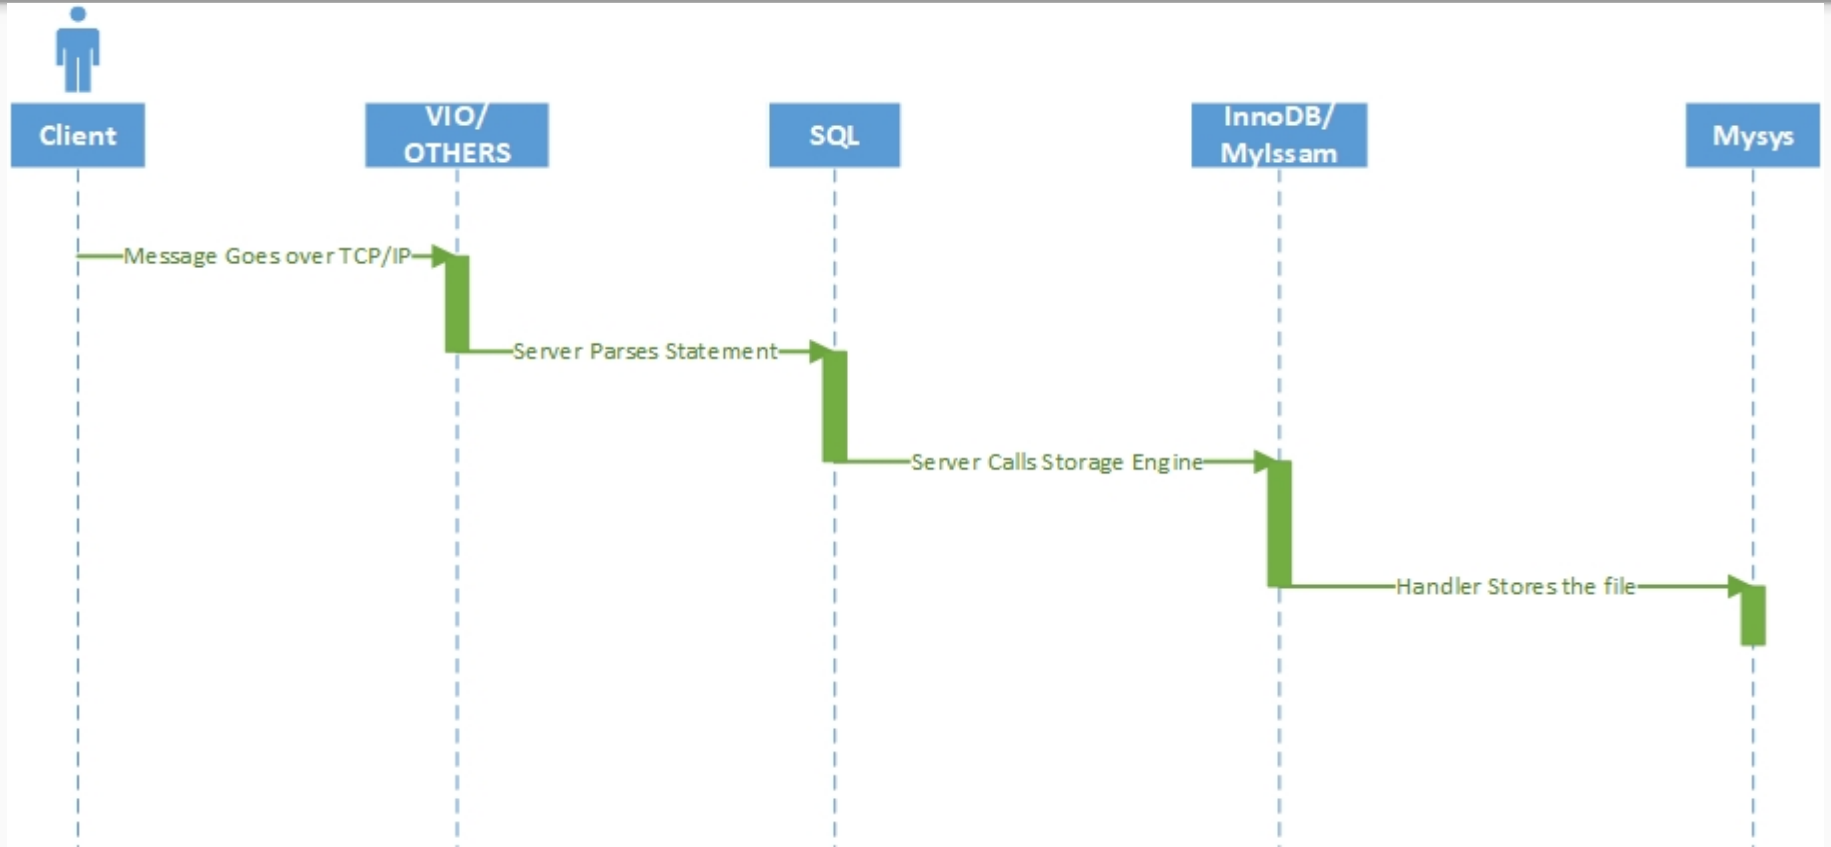








Use Case - INSERT Statement



Lessons Learned

- Retrieving Concrete architecture of a good documented software is possible to achieve with less problems.
- Concrete architecture may contain components that don't exist in the conceptual architecture.
- There are more interactions between components in the concrete architecture than the conceptual.
- Lsedit is a powerful tool but not as user friendly.